

JINMING NIAN

jnian@scu.edu ◇ (530)854-1122 ◇ Santa Clara, CA ◇ [Google Scholar](#) ◇ [Website](#)

SUMMARY

Experienced research assistant and PhD candidate specializing in ML and AI, with a strong background in retrieval, NLP, retrieval augmented generation, and AI reasoning. Skilled in scientific programming with Python/PyTorch, with additional experience in software development and distributed systems. Actively publishing research on retrieval and LLM-related topics.

EDUCATION

Santa Clara University

Ph.D. in Computer Science and Engineering

September 2023 – June 2028 (Expected)

Advisor: [Prof. Yi Fang](#)

M.S. in Computer Science and Engineering

September 2021 – June 2023

Relevant Courses: Machine Learning, Deep Learning, Information Retrieval, Natural Language Processing, Reinforcement Learning

University of California, Davis

September 2017 – June 2021

B.S. in Physics, Minor in Music

WORK EXPERIENCE

Data Scientist Intern, Amazon

June 2025 - September 2025

- Team: Fulfillment Optimization; Organization: Supply Chain Optimization Technology

Research Assistant, Santa Clara University

April 2023 - Present

- Evaluated the reranking effectiveness of LLMs based on their likelihood of generating correct answers for a given query and diverse document contexts, producing a synthetic dataset with quality comparable to human labels for training dense retrievers
- Developed methods using LLMs to generate controlled hallucinations from news articles, creating highly relevant questions that ask for information absent in the document, mimicking RAG scenarios where the retrieved document is unfit for the question
- Analyzed the usefulness of LLM internal attention matrices as features for training a document reranking model
- Developed web-based tools to facilitate human annotation, results aggregation, and inter-annotator agreement calculations

PUBLICATIONS

• W-RAG: Weakly Supervised Dense Retrieval in RAG for Open-domain Question Answering

ACM SIGIR ICTIR 2025 – [paper](#), [code](#)

[Jinming Nian](#), Zhiyuan Peng, Qifan Wang, Yi Fang

• ELOQ: Resources for Enhancing LLM Detection of Out-of-Scope Questions

ACM SIGIR Resource & Reproducibility Track 2025 – [paper](#), [code](#), [dataset](#)

Zhiyuan Peng*, [Jinming Nian](#)*, Alexandre Evfimievski, Yi Fang

• Does Reasoning Introduce Bias? A Study of Social Bias Evaluation and Mitigation in LLM Reasoning

Under Review – [paper](#)

Xuyang Wu*, [Jinming Nian](#)*, Ting-Ruen Wei*, Zhiqiang Tao, Hsin-Tai Wu, Yi Fang

PROJECTS

R1-Zero for Information Retrieval

February 2025 - Present

- Gathered training data: filtered hard queries where both lexical and dense retrieval methods fail; collected irrelevant documents that are consistently retrieved by strong retrievers as hard negative documents
- Redesigned reward modeling based on GRPO for the re-ranking task to encourage self-discovered CoT
- Leveraged Huggingface's Open-R1 as the base framework to train, test, and evaluate on the challenging BRIGHT dataset

Atari Agent (ELEN 552, Reinforcement Learning)

November 2023 - December 2023

- Trained an AI agent to play Pong using Reinforcement Learning methods including a Deep Q-learning Network, Agent memory for experience replay, fixed Q-target, and the epsilon-greedy strategy
- Achieved 89% win-rate against a programmed player with an average winning score of 21:15